

Fraud Detection Dataset Analysis and Machine Learning Classifiers' Comparison for Imbalanced Data

Dimitris Petratos

Abstract

In this study we will compare the application of machine learning classifiers to imbalance data for fraud detection. We will use techniques of oversampling such as SMOTE and we will focus to the influence of them at False Negative Rate, which we consider as the most important metric of rare events' classification such as the one we study.

Imbalance Data

We come across with imbalance data when we have to deal with classification tasks that do not have equal representation from both of the classes. We consider one class to be minor and the other one major class. Cases that could be interpreted as classification of imbalanced data are the one we study here, or others that contain the context of rarity of one of classes such as the occurrence of a rare disease. There have been studied a lot of approaches to problems of imbalance classification. Here we will use oversampling and specifically SMOTE [1] and a special case of that method, Borderline SMOTE [2]. SMOTE creates synthetic data geometrically, producing new registrations between k neighbors of the minor class. Borderline SMOTE does the same by not taking under consideration of the noisy data.

Methodology and Metrics


In the present study, we will use 25/75 test/train split. As evaluation metrics we will measure accuracy, precision, f1-score and recall. What we consider as the most important measure for datasets such as the one we study here, is False Negative Rate (FNR). FNR refers to the cases that an actual rare event is classified as a normal event. FNR is formally defined as $\frac{fn}{fn+tp}$ and our goal is to minimize it. As one can understand, in cases of real time classification, we need a classifier that without any supervision minimize that error. In a real case we would only need human inspection to cases that have been classified as fraud transactions. As we will notice later, when we try to minimize FNR, we will increase False Positive Rate (FPS), formally defined as $FPS = \frac{fp}{fp+tn}$. This is the trade-off that comes up to cases like this one.

We will compare classifiers' performance with and without oversampled train set, holding test set stable.

Data Cleaning and Statistics

For our study we will use Credit Card Fraud Detection dataset from Kaggle. It consists of 284,807 transactions, 492 of which are labeled as fraud. We have 28 features produced from PCA for anonymity reasons and one feature that corresponds to the seconds elapsed between each transaction and the first transaction in the dataset. We also have the feature Amount that corresponds to the amount of the transaction. In the process of data cleaning we found out that the dataset contains 1,081 duplicates that were removed. No null values


were found. After duplicates removal, 473 fraud transactions were left.



So finally we are left with 99.83% of normal transactions and 0.17% of fraud transactions. That indicates the demand of the oversampling that we used.

After train/test split we are left with the same ratio of normal/fraud transactions to train and test set.

We give below the distributions of all the features per class:




We do not observe many differences between the two classes.

Measurements and Results

Logistic Regression

At first we searched for the hyperparameters that we will use for our model. After the following measurements using the test set:



we choose to use l1 penalty with C=1.

For the above model we take the following results:


Accuracy: 0.9991

Precision: 0.8602

F1-Score: 0.7373

Recall: 0.6451

We observe a huge gap between precision and recall.




FNR: 0.3548

FPR: 0.0001

We see that we have almost 35% FNR and that is considered as a lot.

After creating synthetic data with SMOTE, using 5 neighbors, we balance the classes at train set:



and continue implementing our classifier, getting the following results for test set:


Accuracy: 0.9913

Precision: 0.1576

F1-Score: 0.2687

Recall: 0.9112

We take a huge difference for precision and recall, comparing the values we got without oversampling.



FNR: 0.0887

FPR: 0.0085


We take a big drop at the value of FPR and proportionally a big increase at the value of FNR (almost 100 times rise). Finally we use Borderline SMOTE getting the following results:

Accuracy: 0.9973

Precision: 0.3829

F1-Score: 0.5320

Recall: 0.8709



FNR: 0.1290

FPR: 0.0024

For the case of logistic regression we saw that using SMOTE, we take better results regarding the values of FNR, losing from the perspective of precision and FPS. At every case, we have good values for the accuracy.

Borderline SMOTE doesn't seem to affect our results better than the classic implementation of SMOTE, probably because we don't have noise between our classes.

Gaussian Naive Bayes

Using Gaussian Naive Bayes implementation from scikit-learn we don't have to worry about hyperparameter tuning, because there are no extraordinary hyperparameters.


Without oversampling we get the following results:

Accuracy: 0.9923

Precision: 0.1467

F1-Score: 0.2391

Recall: 0.6451



FNR: 0.3548

FPR: 0.0065



We don't get good metrics, as expected.

Below we present our results after SMOTE is used to generate synthetic data: Accuracy: 0.9924

Precision: 0.1595

F1-Score: 0.2650

Recall: 0.7822



FNR: 0.2177

FPR: 0.0072

We don't get better results comparing to the case of logistic regression. Below are the results we get from Borderline SMOTE: Accuracy: 0.9957

Precision: 0.2291

F1-Score: 0.3310

Recall: 0.5967


We choose max_depth=5 and criterion=gini, getting the following results:

Accuracy: 0.9994

Precision: 0.8796

F1-Score: 0.8189

Recall: 0.7661



FNR: 0.2338

FPR: 0.0001


At first, it seems that we get good results. Below are the values after SMOTE oversampling:

Accuracy: 0.9868

Precision: 0.1089

F1-Score: 0.1946


Recall: 0.9112



FNR: 0.4032

FPR: 0.0035

We see that with Gaussian Naive Bayes things go better with SMOTE, but are getting worse when Borderline SMOTE is applied.



FNR: 0.0887

FPR: 0.0130


FNR seems to have been dropped significantly.

Finally the results after Borderline SMOTE are given below:
Accuracy: 0.9965

Precision: 0.3183

F1-Score: 0.4638

Recall: 0.8548



FNR: 0.1451

FPR: 0.0032

We see that with random forest classifier, as measured with all the other classifiers, we get better results with the classic approach of SMOTE, getting a little worse with Borderline SMOTE.

Results and Future Work

We studied 3 classifiers and the influence that synthetic data

has to their performance, since we dealt with imbalance data. We saw that decision tree classifier outperforms logistic regression and Gaussian naive Bayes classifiers using SMOTE technique for synthetic data production. For that reason we considered False Negative Rate as the most important metric measure, since in real cases where there would not be any human interaction with the classifier, we would like to have as few False Negative missclassified documents as we could.

As future work we could try other oversampling methods such as GANs or ADASYN, or even undersampling techniques.

References

1. SMOTE: Synthetic Minority Over-sampling TechniqueNitesh V. Chawla, Kevin W. Bowyer, Lawrence O. Hall, W. Philip Kegelmeyer, 2002
2. Borderline-SMOTE: A New Over-Sampling Method in Imbalanced Data Sets Learning, Hui Han, Wen-Yuan Wang, Bing-Huan Mao, International Conference on Intelligent Computing ICIC 2005: Advances in Intelligent Computing pp 878-887